

Deadlock

Organized By: Vinay Arora
Assistant Professor
CSED, TU

Disclaimer

This is NOT A COPYRIGHT MATERIAL

Content has been taken mainly from the following books:

Operating Systems Concepts By Silberschatz & Galvin,
Operating Systems: Internals and Design Principles By William Stallings

www.os-book.com

www.cs.jhu.edu/~yairamir/cs418/os2/sld001.htm

www.personal.kent.edu/~rmuhamma/OpSystems/os.html

[http://msdn.microsoft.com/en-us/library/ms685096\(VS.85\).aspx](http://msdn.microsoft.com/en-us/library/ms685096(VS.85).aspx)

<http://www.computer.howstuffworks.com/operating-system6.htm>

<http://williamstallings.com/OS/Animations.html>

Etc...

Deadlock – Real Life Scenario



VA.
CSED,TU

System Model

- Resource types R_1, R_2, \dots, R_m

CPU cycles, memory space, I/O devices

- Each resource type R_i has W_i instances.
- Each process utilizes a resource as follows:
 - request
 - use
 - release

Deadlock Characterization

- *Mutual Exclusion*: only one process at a time can use a resource.
- *Hold and Wait*: a process holding at least one resource is waiting to acquire additional resources held by other processes.
- *No Preemption*: a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- *Circular Wait*: there exists a set $\{P_0, P_1, \dots, P_{n-1}\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by

P_2, \dots, P_{n-1} is waiting for a resource that is held by


P_n , and P_0 is waiting for a resource that is held by P_0 .

RAG

- A set of vertices V and a set of edges E .
- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system.
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system.
- request edge – directed edge $P_i \rightarrow R_j$
- assignment edge – directed edge $R_j \rightarrow P_i$

Representation


- Process




- Resource Type with 4 instances




- P_i requests instance of R_j




- P_i is holding an instance of R_j




Example - RAG



RAG with Deadlock



RAG with no Deadlock



Basic Facts


- If graph contains no cycles \Rightarrow no deadlock.
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock.
 - if several instances per resource type, possibility of deadlock.

Deadlock Prevention

- Mutual Exclusion – not required for sharable resources; must hold for non sharable resources.
- Hold and Wait – must guarantee that whenever a process requests a resource, it does not hold any other resources.
 - Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none.
 - Low resource utilization; starvation possible.

- No Preemption –
 - If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released.
 - Preempted resources are added to the list of resources for which the process is waiting.
 - Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.
- Circular Wait – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration.

Safe, Unsafe & Deadlock State



Banker's Algorithm – Safety Algo.

1. Let Work and Finish be vectors of length m and n, respectively.
Initialize Work = Available and Finish[i]=False for $i=0,1,2,\dots,n-1$
2. Find an index i such that both
 Finish[i] == False
 Need_i < Work
 If No such i exists, go to step 4.
3. Work = Work + Allocation_i
 Finish[i] = True
 Go to Step 2.
4. If Finish[i] == true for all i, then the system is in a Safe State.

Banker's Algo. – Resource Request Algo.

1. If $\text{Request}_i \leq \text{Need}_i$, go to step 2. Otherwise, raise an error condition, since the process has exceeded its maximum claim.
2. If $\text{Request}_i \leq \text{Available}$, go to step 3. Otherwise, P_i must wait, since the resources are not available.
3. Have the system pretend to have allocated the requested resources to process P_i by modifying the state as follows:
$$\text{Available} = \text{Available} - \text{Request}_i$$
$$\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i$$
$$\text{Need}_i = \text{Need}_i - \text{Request}_i$$

Problem Statement

| | <u>Allocation</u> | | | <u>Max</u> | | | <u>Available</u> | | |
|----------------|-------------------|---|---|------------|---|---|------------------|---|---|
| | A | B | C | A | B | C | A | B | C |
| P ₀ | 0 | 1 | 0 | 7 | 5 | 3 | 3 | 3 | 2 |
| P ₁ | 2 | 0 | 0 | 3 | 2 | 2 | | | |
| P ₂ | 3 | 0 | 2 | 9 | 0 | 2 | | | |
| P ₃ | 2 | 1 | 1 | 2 | 2 | 2 | | | |
| P ₄ | 0 | 0 | 2 | 4 | 3 | 3 | | | |

Request₁ = (1,0,2)

After Calculating the Matrix Execute for below mentioned Requests

Request₀ = (0,2,0)

Request₄ = (3,3,0)

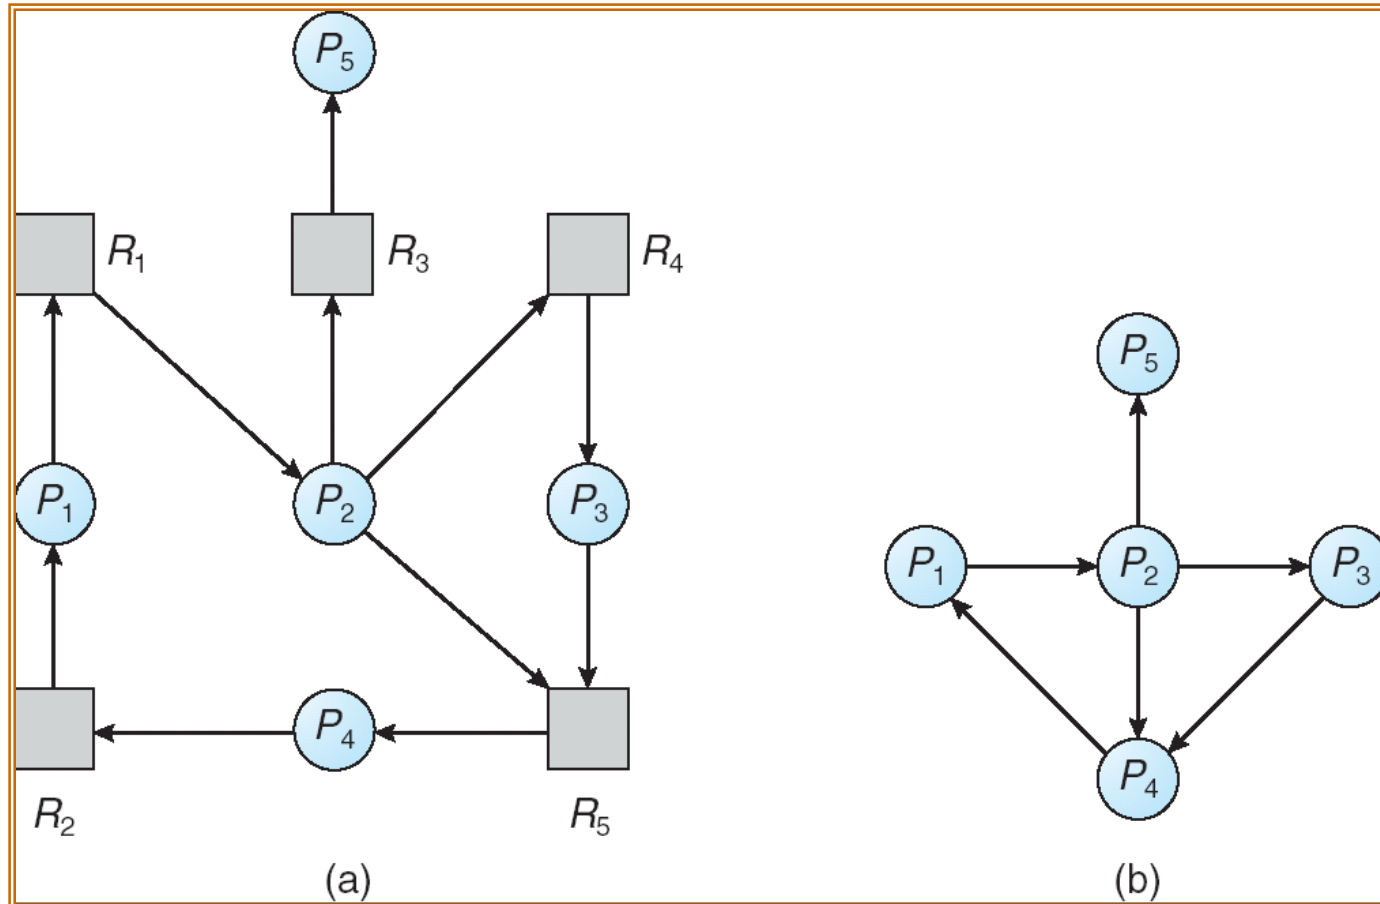
Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme

Single Instance of each Resource Type

- Maintain wait-for graph
- Nodes are processes.
 - $P_i \rightarrow P_j$ if P_i is waiting for P_j .
- Periodically invoke an algorithm that searches for a cycle in the graph.

RAG & Wait For Graph



Several Instances of resource Type

- Available: A vector of length m indicates the number of available resources of each type.
- Allocation: An $n \times m$ matrix defines the number of resources of each type currently allocated to each process.
- Request: An $n \times m$ matrix indicates the current request of each process. If Request $[ij] = k$, then process P_i is requesting k more instances of resource type. R_j .

Deadlock Detection

1. Let Work and Finish be vectors of length m and n, respectively.
Initialize Work = Available. For $i=0,1,2,\dots,n-1$, if Allocation_i Not equal to ZERO, then $\text{Finish}[i] = \text{False}$; Otherwise, $\text{Finish}[i] = \text{True}$
2. Find an index i such that both
 $\text{Finish}[i] == \text{False}$
 $\text{Request}_i \leq \text{Work}$
If No such i exists, go to step 4.
3. $\text{Work} = \text{Work} + \text{Allocation}_i$
 $\text{Finish}[i] = \text{True}$
Go to Step 2.
4. If $\text{Finish}[i] == \text{false}$ for some i, then the system is in Deadlock State.

Problem Statement

| | <u>Allocation</u> | | | <u>Request</u> | | | <u>Available</u> | | |
|----------------|-------------------|---|---|----------------|---|---|------------------|---|---|
| | A | B | C | A | B | C | A | B | C |
| P ₀ | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| P ₁ | 2 | 0 | 0 | 2 | 0 | 2 | | | |
| P ₂ | 3 | 0 | 3 | 0 | 0 | 0 | | | |
| P ₃ | 2 | 1 | 1 | 1 | 0 | 0 | | | |
| P ₄ | 0 | 0 | 2 | 0 | 0 | 2 | | | |

| <u>Request</u> | | |
|----------------|---|---|
| A | B | C |
| 0 | 0 | 0 |
| 2 | 0 | 2 |
| 0 | 0 | 1 |
| 1 | 0 | 0 |
| 0 | 0 | 2 |

Recovery from Deadlock – Process Termination

- Abort all deadlocked processes.
- Abort one process at a time until the deadlock cycle is eliminated.
- In which order should we choose to abort?
 - Priority of the process.
 - How long process has computed, and how much longer to completion.
 - Resources the process has used.
 - Resources process needs to complete.
 - How many processes will need to be terminated.

Recovery from Deadlock – Resource Preemption

- Selecting a victim – minimize cost.
- Rollback – return to some safe state, restart process for that state.
- Starvation – same process may always be picked as victim, include number of rollback in cost factor.



Thnx...